

Q1. Describe how Linked List is implemented in Pintos.

In Pintos, a list contains two `list_elem` which are head and tail respectively while each linked list element contains pointers to its previous and next elements and thus a linked list does not use dynamic memory allocation. Particularly, a `list_elem` can only be part of one list at a time since its previous and next `list_elem` are confirmed.

To construct a list with specific type, for example Integer, we use `int_list_elem` which contains an integer and a `list_elem`.

For construction, we declare and initialize a list and use `list_push_back()` method to push a `list_elem`, instead of `int_list_elem`, into this list. Method `list_entry()` can get the integer value by using the offset of `int_list_elem` and `list_elem`.

Methods for sort and insert are also available with provided list element comparison function.

Q2. How to use integer to simulate real number calculation(Suggested examples)?

We can use fixed-point number representation to use integer to simulate real number.

For example, for a 32-bit integer, we divide it into 3 parts which are one sign bit, 15 bits ahead of the decimal point and the remaining rightmost 16 bits for fraction part. In this case, a real number times 2^{16} will be used in Pintos to represent this number and the maximum integer representation 32-bit 1 in Pintos becomes $(2^{31} - 1)/(2^{16})$ for real number value.

Many operations, like addition and subtraction, on fixed-point numbers are straightforward. While for multiplication and division we have to divide or multiply an extra 2^{16} to rebalance the results and sometimes use 64-bit operation in case of overflow.

Suppose x, y are fixed-point representation. For addition and subtraction, it is $(x+y)$ and $(x-y)$ while for multiplication and division, it is $((\text{int64_t})x*y/2^{16})$ and $((\text{int64_t})x*2^{16})/y$.

Q3. Why can multi-level feedback scheduling avoid starvation?

Thread priority is recalculated once every fourth clock tick according to $\text{priority} = \text{PRI_MAX} - (\text{recent_cpu}/4) - (\text{nice} \times 2)$.

The scheduler has 64, from 0 to 63, priorities with 63 representing the `PRI_MAX` which has the highest priority.

Integer "nice" denotes how nice a thread is compared to other threads with a positive nice value cause a thread decreasing its priority and a negative one receiving more CPU time.

"recent_cpu" is an estimate of the CPU time the thread has used recently. It can preventing

starvation since a thread that has not received any CPU time recently will have a `recent_cpu` of 0 and thus it will not decrease its priority because of `recent_cpu` and receives CPU time soon while threads that have recently been scheduled on the CPU will have a lower priority the next time the scheduler picks a thread to run.